

[Return to Classroom](#)

Data Modeling with Cassandra

REVIEW

HISTORY

Requires Changes

1 specification requires changes

Hi there,

This is an excellent submission. Most of the suggested changes are now in place. Congratulations on coming so far. The

submission shows that you have understood the concepts of dimensional modeling in Cassandra pretty well.

Some of the key points you should feel really confident now as you did well in the submission:



Selection of datatypes.



Creating separate tables for each query.



Using appropriate fields in SELECT queries.



Maintaining order of COMPOSITE KEYS to CREATE and INSERT statements.

However, you need to make a subtle but critical change, that the last reviewer missed to point out



Correcting the COMPOSITE KEY of the third table, you should not use session data there.

I have tried to explain the required changes in detail and how to make those changes. Kindly go through them. It would be hardly a few minutes task for you as everything else is in order. I hope you will take this constructively and consider it as an opportunity to learn and grow. 😊



We hope the lessons gave a very good base but there is a lot to explore. You can take a look at this free course on [Apache Cassandra](#) from DataStax itself if you want to hone your skills on Cassandra Data Modeling.

Wish you all the very best for your next submission and all your endeavors. 👍

ETL Pipeline Processing


Student creates `event_data_new.csv` file.

Perfect!! The ETL process creates the `event_data_new.csv` from the log files in the staging area.
You have followed the instructions properly and there are 6821 records as expected.

Student uses the appropriate datatype within the `CREATE` statement.

Good work with the datatypes selection for the fields. The very basic importance of assigning correct datatype is when we need to do data wrangling/aggregation. What does that mean? For example, if you are asked to sum up the total length of songs a user listened to in a particular session, having the length's datatype as the float will help to get the end result perfectly.


- `song`, `first_name`, `last_name` : TEXT, ✓
- `session_id`, `item_in_session`, `userId` : INT ✓
- `length` : DECIMAL ✓


 While data modeling you may like to keep this [official documentation](#) on datatypes handy for future Cassandra project when assigning the correct data types.

Data Modeling

Student creates the correct Apache Cassandra tables for each of the three queries. The `CREATE TABLE` statement should include the appropriate table.

You have adhered to one table per query rule of Apache Cassandra. No two queries are answered from the same table. This is important as this decides the efficiency of query retrieval.

 It is understandable that keeping similar data will eat a lot of space. *But if you compare the cost of storage to CPU, the former will be always preferred.*

 This [documentation](#) clearly describes the fundamental modeling concept of Cassandra, why data duplication in separate

tables are allowed.

Student demonstrates good understanding of data modeling by generating correct SELECT statements to generate the result being asked for in the question.

The SELECT statement should NOT use `ALLOW FILTERING` to generate the results.

Awesome

The specific field names are used to retrieve data. **This is also a very important point when we are accessing a very huge database and doing `SELECT *` can be very costly in terms of reading.** As I mentioned above which is a basic principle of Cassandra data modeling, *Prefer storage cost over CPU cost or read/write cost.*, so we do not really mind keeping duplicate data, but while reading we must be alert not to increase IO operations.

✓ First query

```
"select artist_name, song_name, song_length from songs_in_session where session_id = 338 and  
item_in_session = 4"
```

✓ Second query

```
"select item_in_session, artist_name, song_name, user_first_name, user_last_name from artist_s  
ong_in_session where session_id = 182 and user_id = 10"
```

✓ Third query

```
"select user_first_name, user_last_name from user_by_song where song_name = 'All Hands Against His Own'"
```

Student should use table names that reflect the query and the result it will generate. Table names should include alphanumeric characters and underscores, and table names must start with a letter.

Great work here!! The table names reflect the query. **Cassandra is a query first database where data models/tables are created based on the query it is going to answer** (hence we create different tables for different queries). To help the stakeholders to quickly understand what is going to be the use of any specific table, **a name which reflects the query will be more helpful.**

- `songs_in_session` ✓
- `artist_song_in_session` ✓
- `user_by_song` ✓

The sequence in which columns appear should reflect how the data is partitioned and the order of the data within the partitions.

Perfect

You have done a splendid job here by maintaining the sequence of CREATE and INSERT to the order of COMPOSITE PRIMARY KEY and CLUSTERING COLUMNS. This is one of the key takeaways from this project and you did it perfectly. **The data should be inserted and retrieved in the same order as to how the COMPOSITE PRIMARY KEY is set up.** This is important because

Apache Cassandra is a partition row store, which means the partition key determines where any particular row is stored and on which node.

 Here is a [DataStax documentation](#) for you further understanding. and an amazing [article](#) on Cassandra Data Modelling.

PRIMARY KEYS

The combination of the **PARTITION KEY** alone or with the addition of **CLUSTERING COLUMNS** should be used appropriately to uniquely identify each row.

- ✓ The COMPOSITE KEY combination for the first table is in place. The `session_id` and `item_in_session` perfectly filters the data as per the query we are answering from the table.
- ✓ You have explicitly defined `item_in_session` as **CLUSTERING COLUMN** for the second table to fulfill the sorting requirement.

✗ It Seems, the last reviewer has missed this or overlooked a small problem here. Apologies for that 😊

For the third query to answer you are using `session_id`, `item_in_session` in your table, and as a part of your COMPOSITE KEY. Let us understand first what is expected: We need to get a particular song and the users who listened to it. That means **we need to have a combination of PRIMARY KEY which uniquely separates each song and user combination**. But in your current implementation, the records will be separated based on different sessions. i.e if a user listens to the song multiple times at different sessions all of the records will be stored. But we need only a record of `user+song` combo. Hence, I would suggest that you remove the `session_id`, `item_in_session` from the table and COMPOSITE KEY and use `user_id` which will help to uniquely identify each record as desired.

Presentation

The notebooks should include a description of the query the data is modeled after.

EXCELLENT

This rubric is to evaluate the project presentation in a professional and polished manner.

Amazing work with the notebook presentation. The Queries are divided into subtasks. Each task has proper headings and descriptions of schema design. 👍

Code should be organized well into the different queries. Any in-line comments that were clearly part of the project instructions should be removed so the notebook provides a professional look.

The TODOs and inline comments are removed. Consider you are presenting this project as this notebook to an employer. We do not need to have all the questions and instructions to be in there. It is more professional when separate them as tasks and mention what we have done and why we have done it. You have done it perfectly. Great job 👍

 RESUBMIT

 [DOWNLOAD PROJECT](#)

RETURN TO PATH

Rate this review

START
